

OceanSaksham: An Open-Source Architecture for Credibility-Aware Triage and Audit Logging in Crowdsourced Coastal Hazard Management

Suraj Gore*, Saachi Sharma†

*Department of Computer Science and Engineering, Ocean Data Systems Lab

†Collaborative Study on Ocean Information Services and Coastal Resilience

Email: {suraj.gore, saachi.sharma}@oceansaksham.gov.in

Abstract—Crowdsourced citizen reporting is an important mechanism for gathering hyper-local situational awareness during coastal emergencies such as storm surges, high wave breaches, tidal inundation, and shoreline erosion. However, emergency operation centres and disaster management agencies frequently encounter operational bottlenecks caused by unverified, noisy, duplicate, or uncalibrated submissions. In this paper, we present the system design, multi-signal feature architecture, and reproducible synthetic benchmark evaluation of OceanSaksham, an open-source platform for coastal hazard reporting featuring an explainable, human-in-the-loop credibility triage layer and tamper-evident audit logging.

The proposed triage engine formulates 19 features across seven modalities: geospatial GPS precision, camera EXIF metadata cross-validation, temporal submission latency, multilingual Indic keyword alignment (across nine coastal Indian languages), duplicate perceptual media hashing, spatial-temporal neighborhood corroboration, and real-time physical sea-state telemetry via the Open-Meteo Marine API. To evaluate the triage mechanics under controlled conditions, we formulate a 5-seed parameterized synthetic benchmark ($N = 1,000$ per seed) evaluated under a strict region-held-out spatial split (training on West Coast regions, evaluating on East Coast regions). Across five independent dataset generations, regularized logistic regression achieves $F_1 = 0.946 \pm 0.004$ (AUC-ROC = 0.988 ± 0.001), random forest achieves $F_1 = 0.938 \pm 0.009$ (AUC-ROC = 0.987 ± 0.002), and gradient-boosted decision trees achieve $F_1 = 0.934 \pm 0.012$ (AUC-ROC = 0.986 ± 0.003). We analyze the benchmark’s properties, noting that rule-based heuristics perform identically to ML models on synthetic data, detail the complete system pipeline, and provide privacy design considerations informed by India’s DPDP Act 2023.

Index Terms—Crisis Informatics, Crowdsourcing, Coastal Hazard Management, Credibility Scoring, System Architecture, Explainable AI, Human-in-the-Loop, Audit Logging.

I. Introduction

Coastal communities across India’s 7,516 km shoreline face continuous threats from severe oceanographic and meteorological hazards, including cyclonic storm surges, extreme swell waves, tidal floods, and rapid shoreline erosion [1], [2]. While institutional remote sensing networks and coastal radar stations operated by agencies such as the Indian National Centre for Ocean Information Services (INCOIS) provide regional forecasts, hyper-local impact assessments (such as breached sea walls, submerged village

roads, or fishing harbour distress) rely on participatory crowdsourcing [3], [4].

However, raw citizen submissions present severe trust and filtering challenges for emergency dispatchers [5], [6]. During acute marine emergencies, dispatch queues become congested with unverified reports, recycled duplicate images from past disasters [7], inaccurate GPS readings, and well-intentioned false alarms. Furthermore, in multilingual regions spanning official languages such as Hindi, Marathi, Tamil, Telugu, Bengali, Gujarati, Malayalam, Odia, and English, uniform manual text triage is slow and error-prone.

To address these challenges, this paper presents the engineering architecture, multi-signal feature formulation, and reproducible synthetic evaluation of OceanSaksham, focusing on four core contributions:

- 1) **Open-Source Full-Stack Architecture:** An end-to-end platform featuring token-bucket rate limiting, Server-Sent Events (SSE) broadcasting, spatial duplicate detection ($< 2\text{km}, < 3\text{h}$), tamper-evident audit logging of official decisions, and a multi-state SOS alert workflow.
- 2) **Multi-Signal Feature Formulation:** A 19-dimensional feature representation spanning seven modalities: geospatial GPS precision, photo EXIF cross-validation, temporal urgency, multilingual Indic hazard keyword alignment, duplicate media hashing, spatial-temporal neighborhood corroboration, and physical sea-state validation via Open-Meteo marine telemetry.
- 3) **Explainable Human-in-the-Loop Console:** A dispatch dashboard that ranks incoming reports by credibility while exposing top positive and negative signal contributions, preserving complete operator authority while logging official overrides for model feedback.
- 4) **Parameterized Synthetic Benchmark & Baseline Harness:** A multi-seed synthetic evaluation harness and region-held-out spatial testbed across 14 coastal regions for systematic baseline benchmarking and future evaluation with real pilot data.

II. Related Work & Platform Comparison

A. Crisis Informatics & Participatory Sensing

Crowdsourcing platforms such as Ushahidi demonstrated the utility of citizen crisis mapping during the 2010 Haiti earthquake [3]. In geophysics, the USGS “Did You Feel It?” (DYFI) system established citizen reports as a valuable complement to physical seismograph arrays [8]. Sakaki et al. [4] utilized social media streams as real-time event sensors for earthquake detection. However, raw social streams often lack structured telemetry, location verification, and spam resistance [9].

B. Credibility Assessment & Disaster NLP

Information credibility classification has traditionally leveraged propagation graphs and linguistic sentiment [5]. Gupta et al. [7] demonstrated that over 86% of disaster images during Hurricane Sandy were recycled fakes, emphasizing the necessity of media hashing and EXIF validation. Recent benchmarks such as HumAID [10] and CrisisNLP [6] have expanded text filtering, while IndicNLP [11] advanced Indic text processing. Our work builds upon these principles by embedding real-time physical marine sensor telemetry (H_s , wind speed) directly into the triage decision space.

C. Platform Comparison Matrix

Table I benchmarks OceanSaksham against representative crisis crowdsourcing platforms.

III. System Architecture & Incident Pipeline

A. System Architecture

As illustrated in Fig. 1, the platform comprises three tiers:

- **Citizen Web Client:** Mobile-responsive interface built with React 18 and Vite, supporting Web App Manifest installation, 9 Indic languages, camera EXIF extraction, and GPS capture.
- **Backend REST & SSE Service:** Node.js/Express service with SQLite storage in development (with spatial Haversine queries, designed for migration to PostgreSQL/PostGIS in production), token-bucket rate limiting (max 30 req/min), spatial duplicate checking, and Server-Sent Events broadcasting.
- **Official Dispatch Console:** Command dashboard displaying ranked reports, explainable signal attribution badges, SOS alert monitoring, and danger zone hotspot management.

B. End-to-End Incident Pipeline Walkthrough

The system operates through an end-to-end automated pipeline:

- 1) **Citizen Ingestion:** A citizen captures a photo and submits a hazard report with GPS coordinates and localized description.
- 2) **Duplicate Screening:** The backend computes the image SHA-256 hash and checks against existing




Fig. 1. Three-tier architecture of the OceanSaksham platform.

reports within $\Delta r \leq 2.0$ km and $\Delta t \leq 3.0$ h. If matched, the report is linked to the primary incident to prevent queue congestion.

- 3) **Environmental Grounding:** The backend queries the Open-Meteo Marine Weather API [12] for offshore wave height (H_s) and wind speed at the nearest maritime grid point.
- 4) **Multi-Signal Scoring:** The 19-dimensional feature vector is evaluated, calculating the composite credibility score along with top positive and negative signal attributions.
- 5) **Operator Decision & Audit Log:** Official operators review the ranked queue on the dispatch console. Any status transition (VERIFIED, REJECTED) generates an immutable record in `audit_logs` capturing operator ID, timestamps, old/new states, and override rationale.

IV. Credibility Triage Feature Formulation

We formulate a 19-dimensional feature vector $\mathbf{x} \in \mathbb{R}^{19}$ across seven groups:

1. **Location Group** ($\mathbf{f}_{\text{loc}} \in \mathbb{R}^4$):
 - $x_1 = \ln(1 + \text{accuracy}_{\text{GPS}})$: Log-normalised GPS error radius.
 - $x_2 = d_{\text{Haversine}}((\text{lat}_{\text{GPS}}, \text{lng}_{\text{GPS}}), (\text{lat}_{\text{EXIF}}, \text{lng}_{\text{EXIF}}))$: Spatial discrepancy between reported location and camera metadata.
 - $x_3 \in \{0, 1\}$: Binary indicator of valid EXIF GPS presence.
 - x_4 : Minimum distance to canonical coastline anchor points.
2. **Time Group** ($\mathbf{f}_{\text{time}} \in \mathbb{R}^3$):
 - $x_5 = \ln(1 + |t_{\text{server}} - t_{\text{device}}|/60)$: Submission latency in minutes.
 - $x_6 = \sin(2\pi h/24), x_7 = \cos(2\pi h/24)$: Diurnal solar cycle encoding.
3. **Text Group** ($\mathbf{f}_{\text{text}} \in \mathbb{R}^3$):
 - $x_8 = \ln(1 + \text{len}(\text{desc}))$: Description character volume.
 - $x_9 \in \{0, 1\}$: Keyword match against localized hazard terminology across Indic lexicons.
 - $x_{10} \in \{0, 1\}$: Language indicator (English vs. vernacular).

TABLE I
Feature Comparison of Crisis Crowdsourcing and Ocean Information Platforms

Capability / Feature	Ushahidi [3]	USGS DYFI [8]	INCOIS SAMUDRA [1]	Sahana Eden	OceanSaksham (Ours)
Primary Domain	General Crisis	Earthquakes	Ocean Forecasts	Relief Logistics	Coastal Hazards
Multilingual Indic UI	Partial (Plugins)	No (EN/ES)	Yes (Govt. Lang.)	Partial	Yes (9 Coastal Languages)
Camera EXIF Cross-Check	No	No	No	No	Yes ($\Delta r_{\text{GPS-EXIF}}$)
Spatial Deduplication	Manual	Grid-Aggregated	No	Partial	Automated ($< 2\text{km}, < 3\text{h}$)
Marine Sensor Triage	No	No (Physical Array)	Yes (Broadcast only)	No	Yes (Open-Meteo Integration)
Explainable Credibility Score	No (Manual)	Rule Intensity (CDI)	No	No	Yes (7 Feature Groups)
Tamper-Evident Audit Log	No	No	Internal Logs	Role-based	Yes (Immutable Decision Trail)
Multi-State SOS Workflow	No	No	No	Incident Tracking	Yes (State Machine + SSE)

4. Media Group ($\mathbf{f}_{\text{media}} \in \mathbb{R}^2$):

- $x_{11} \in \{0, 1\}$: Attached media presence.
- $x_{12} \in \{0, 1\}$: Duplicate SHA-256 media hash match.

5. Corroboration Group ($\mathbf{f}_{\text{corrob}} \in \mathbb{R}^2$):

- $x_{13} = \sum_{j \neq i} \mathbb{I}[d(r_i, r_j) \leq 5\text{km} \wedge |t_i - t_j| \leq 6\text{h}]$: Independent nearby report count.
- x_{14} : Hazard-type agreement ratio among neighbors.

6. Reporter Reputation ($\mathbf{f}_{\text{rep}} \in \mathbb{R}^2$):

- x_{15} : Historical verification rate.
- $x_{16} = \ln(1 + N_{\text{total}})$: Historical report volume.

7. External Marine Weather ($\mathbf{f}_{\text{ext}} \in \mathbb{R}^3$):

- x_{17} : Significant wave height (H_s in meters) retrieved from Open-Meteo Marine API.
- x_{18} : Coastal wind speed (km/h).
- $x_{19} = H_s \times \text{Severity}_{\text{reported}}$: Sea-state consistency interaction term.

V. Experimental Methodology & Benchmark Suite

A. Synthetic Dataset Generation & Threat Models

Due to the absence of public, verified crowdsourced coastal disaster datasets containing multi-modal telemetry and EXIF metadata, we developed a parameterized synthetic benchmark suite across 14 coastal regions (7 West Coast: Mumbai, South Gujarat, North Gujarat, Goa, Coastal Karnataka, North Kerala, South Kerala; 7 East Coast: Chennai, South Tamil Nadu, Andhra Central, Visakhapatnam, Odisha South, Odisha North, West Bengal).

Each synthetic seed ($N = 1,000$) simulates realistic threat conditions:

- Adverse Weather GPS Degradation: Genuine storm surge submissions with degraded GPS accuracy ($\pm 50\text{--}150\text{m}$).
- Messaging App Metadata Stripping: 50% of authentic images simulated without EXIF tags, reflecting typical messaging compression.
- Adversarial Hoaxes: High-precision inland reports contradicted by calm sea telemetry ($H_s < 1.3\text{m}$) and coastal distance.
- Region-Held-Out Spatial Partition: 578 West Coast reports for training; 422 East Coast reports reserved exclusively for testing.

VI. Results & Evaluation

A. Classifier Performance (5-Seed Region-Held-Out Evaluation)

To ensure statistical reproducibility, we generated 5 independent datasets across seeds {42, 123, 456, 789, 2024} and evaluated models on the held-out East Coast partition ($N = 422$). Table II reports the Mean $\pm$ Standard Deviation across all 5 seeds.

All three models perform within standard deviation overlap. This demonstrates that standard linear and ensemble classifiers reliably learn the multi-signal representation without overfitting to specific regional partitions.

B. Queue Ranking Metrics

Table III evaluates dispatch queue prioritization across the 5 benchmark seeds.

Notably, the hand-crafted rule heuristic matches the ML model ($P@k = 1.00$, NDCG@10 = 1.000). We interpret this not as evidence that rule systems are inherently superior in production, but as an expected property of the synthetic data generator, which models specific generative rules. The system’s primary contribution is providing an explainable architecture and baseline harness that supports both rule-based heuristics and learned models interchangeably.

C. Feature Group Ablation

Table IV details the effect of systematically removing each feature group across the 5 evaluation seeds.

In this synthetic setup, omitting external marine weather leads to the largest drop in AUC ($\Delta\text{AUC} = -0.015 \pm 0.006$), reflecting the generator’s inclusion of calm-water hoax scenarios.

VII. Limitations & Ethical Considerations

Limitations:

- 1) Synthetic Benchmark Constraints: All quantitative results in this paper are derived from a synthetic benchmark testbed. While the generator incorporates domain threat models, performance on real-world disaster streams must be validated in future field deployments.
- 2) Prevalence & Class Skew: The benchmark exhibits an active hazard prevalence of $61.9\% \pm 3.7\%$, which

TABLE II
Classifier Performance on Region-Held-Out Spatial Test Set (Mean $\pm$ Std across 5 Independent Seeds, $N = 422$)

Model	Precision	Recall	F_1 -Score	AUC-ROC	AUC-PR
Logistic Regression (L2)	0.952 $\pm$ 0.013	0.941 $\pm$ 0.011	0.946 $\pm$ 0.004	0.988 $\pm$ 0.001	0.993 $\pm$ 0.001
Random Forest ($n = 100$)	0.940 $\pm$ 0.021	0.937 $\pm$ 0.025	0.938 $\pm$ 0.009	0.987 $\pm$ 0.002	0.992 $\pm$ 0.002
GBDT (Gradient Boosting)	0.932 $\pm$ 0.014	0.937 $\pm$ 0.017	0.934 $\pm$ 0.012	0.986 $\pm$ 0.003	0.991 $\pm$ 0.003

TABLE III
Queue Ranking Metrics across 5 Seeds (Mean $\pm$ Std)

Strategy	P@5	P@10	P@20	NDCG@10
FIFO Queue	0.40 $\pm$ 0.28	0.34 $\pm$ 0.26	0.32 $\pm$ 0.27	0.381 $\pm$ 0.240
Severity Heuristic	0.68 $\pm$ 0.10	0.70 $\pm$ 0.14	0.65 $\pm$ 0.11	0.708 $\pm$ 0.154
Rule-Based Score	1.00 $\pm$ 0.00	1.00 $\pm$ 0.00	1.00 $\pm$ 0.00	1.000 $\pm$ 0.000
ML Triage (GBDT)	1.00 $\pm$ 0.00	1.00 $\pm$ 0.00	1.00 $\pm$ 0.00	1.000 $\pm$ 0.000

TABLE IV
Feature Group Ablation Study across 5 Seeds (GBDT)

Configuration	AUC-ROC	F_1 -Score	Δ AUC
Full 7-Group Vector	0.986 $\pm$ 0.003	0.934 $\pm$ 0.012	0.000
Without Location	0.985 $\pm$ 0.003	0.931 $\pm$ 0.011	-0.001 $\pm$ 0.001
Without Time	0.985 $\pm$ 0.003	0.933 $\pm$ 0.013	-0.001 $\pm$ 0.001
Without Text	0.986 $\pm$ 0.003	0.935 $\pm$ 0.011	+0.000 $\pm$ 0.001
Without Media	0.985 $\pm$ 0.003	0.932 $\pm$ 0.012	-0.001 $\pm$ 0.001
Without Corroboration	0.985 $\pm$ 0.003	0.933 $\pm$ 0.012	-0.001 $\pm$ 0.001
Without Reporter	0.986 $\pm$ 0.003	0.934 $\pm$ 0.011	+0.000 $\pm$ 0.001
Without Marine Weather	0.971 $\pm$ 0.006	0.906 $\pm$ 0.018	-0.015 $\pm$ 0.006

is substantially higher than typical non-crisis operational streams where genuine emergencies may represent a small minority of incoming reports.

- 3) Marine API Spatial Resolution: Open-Meteo marine models interpolate offshore open-water grid cells; localized coastal structures (harbour breakwaters, shallow reef estuaries) may exhibit hydrodynamic variations not captured by open-ocean wave models.
- 4) SOS Emergency Dispatch: The SOS emergency module is a prototype interface workflow and does not interface directly with national emergency dispatch services (such as India’s 112 system).

Privacy Design Considerations: The platform incorporates privacy protections informed by India’s Digital Personal Data Protection (DPDP) Act 2023 [13]:

- Purpose Limitation & Spatial Coarsening: Public citizen map views coarsen location coordinates to 2 decimal places ($\approx$ 1.1 km precision) to prevent exposing private residences.
- Pseudonymization: User identifiers are tokenized using salted SHA-256 hashes in public-facing SSE event broadcasts.
- Immutable Audit Trail: Official triage actions are stored in an append-only log to ensure accountability and operator oversight.

VIII. Conclusion & Future Work

We presented OceanSaksham, an open-source architecture for credibility-aware triage and audit logging in

crowdsourced coastal hazard management. By coupling multi-signal feature engineering with real-time oceanographic validation, the platform offers an explainable, human-in-the-loop dispatch environment. Future work includes field piloting with regional disaster management authorities and INCOIS, collecting real citizen report streams, and exploring on-device edge ML for low-connectivity coastal fishing operations.

Acknowledgment

This research was developed as part of the OceanSaksham coastal resilience initiative. Oceanographic weather data is accessed via the Open-Meteo Marine API under Open Database License (ODbL) terms.

References

- [1] T. S. Kumar, C. Sheno, R. S. Mahendra, and B. A. Kumar, “Coastal vulnerability assessment for Orissa state, east coast of India,” *Natural Hazards*, vol. 64, no. 3, pp. 2037–2052, 2012.
- [2] N. Balakrishnan, S. K. Dash, and U. C. Mohanty, “Simulation of very severe cyclonic storms over the Bay of Bengal using mesoscale models,” *Atmospheric Research*, vol. 209, pp. 132–145, 2018.
- [3] O. Okolloh, “Ushahidi, or ‘testimony’: Web 2.0 tools for crowdsourcing crisis information,” *Participatory Learning and Action*, vol. 59, no. 1, pp. 65–70, 2009.
- [4] T. Sakaki, M. Okazaki, and Y. Matsuo, “Earthquake shakes Twitter users: Real-time event detection by social sensors,” *Proceedings of the 19th International Conference on World Wide Web (WWW)*, pp. 851–860, 2010.
- [5] C. Castillo, M. Mendoza, and B. Poblete, “Information credibility on Twitter,” *Proceedings of the 20th International Conference on World Wide Web (WWW)*, pp. 675–684, 2011.
- [6] M. Imran, C. Castillo, F. Diaz, and S. Vieweg, “Processing social media messages in mass emergency: A survey,” in *ACM Computing Surveys (CSUR)*, vol. 47, no. 4, 2015, pp. 1–38.
- [7] A. Gupta, H. Lamba, P. Kumaraguru, and A. Joshi, “Faking Sandy: Characterizing and identifying fake images on Twitter during Hurricane Sandy,” *Proceedings of the 22nd International Conference on World Wide Web (WWW)*, pp. 729–736, 2013.
- [8] D. J. Wald, V. Quitoriano, T. H. Heaton, and H. Kanamori, “Relationships between peak ground acceleration, peak ground velocity, and modified Mercalli intensity in California,” *Earthquake Spectra*, vol. 15, no. 3, pp. 557–564, 1999.
- [9] C. Reuter, A. L. Hughes, and M.-A. Kaufhold, “Social media in crisis management: An evaluation and analysis of crisis informatics research,” *International Journal of Information Management*, vol. 39, pp. 198–211, 2018.
- [10] F. Alam, U. Qazi, M. Imran, and F. Ofli, “Humaid: Multimodal human-annotated disasters dataset with images and texts,” in *Proceedings of the International AAAI Conference on Web and Social Media (ICWSM)*, vol. 15, 2021, pp. 933–943.
- [11] D. Kakwani, A. Kunchukuttan, S. Golla, G. N.C., A. Bhattacharyya, M. M. Khapra, and P. Kumar, “Indicnlp suite: Monolingual corpus, language models and pre-trained multilingual language models for 22 Indic languages,” in *Findings of EMNLP*, 2020, pp. 4948–4961.

- [12] Open-Meteo, “Open-meteo marine weather forecast API,” <https://open-meteo.com/en/docs/marine-weather-api>, 2024.
- [13] Ministry of Law and Justice, Government of India, “The digital personal data protection act, 2023 (No. 22 of 2023),” 2023, the Gazette of India.